Comprehensive Survey of Attention Mechanisms in Encoder-Only Language Models

Architectures, Efficiencies, and Interpretability

Cristian Leo Daniel Begimher

March 2026

Abstract

The introduction of the Transformer architecture in 2017 catalyzed a fundamental paradigm shift in natural language processing [19]. While decoder-only autoregressive models have come to dominate generative AI, encoder-only bidirectional models historically establish and maintain state-of-the-art results in natural language understanding, dense information retrieval, and sequence classification. This paper presents an exhaustive survey of the self-attention mechanism within encoder-only models. We analyze its mathematical foundations, trace its architectural evolution from the original BERT [7] through RoBERTa [13], DeBERTa [8], and ModernBERT [22], evaluate efficiency enhancements including sparse attention [25], linear kernelized attention, and hardware-aware FlashAttention [6], and dissect the ongoing theoretical debates surrounding attention-based interpretability. We further survey hybrid architectures that interleave state-space models with self-attention, and discuss the practical limits and deployment considerations of encoder attention in long-context regimes.

Contents

1 Introduction to the Encoder-Only Paradigm **3**
1.1 Scope and Contributions . . . 3

2 Mathematical Foundations of Bidirectional Self-Attention **3**
2.1 Scaled Dot-Product Attention . . . 3
2.2 Multi-Head Attention . . . 4
2.3 Computational Complexity . . . 4

3 Architectural Evolution of Foundational Encoders **4**
3.1 BERT: Bidirectional Pre-Training at Scale . . . 4
3.2 RoBERTa: Optimizing the Training Recipe . . . 5
3.3 ALBERT: Parameter-Efficient Factorization . . . 5
3.4 ELECTRA: Sampling-Efficient Discriminative Pre-Training . . . 5

4 Disentangling Context and Position: DeBERTa **6**
4.1 Disentangled Attention Formulation . . . 6

5 Taming the Quadratic Bottleneck **6**
5.1 Sparse and Windowed Attention . . . 6
5.2 Linear and Kernelized Attention . . . 7
5.3 Summary: Complexity Comparison . . . 7

6 Hardware-Aware Optimizations **7**
6.1 The I/O Bottleneck . . . 8
6.2 FlashAttention . . . 8
6.3 FlexAttention . . . 8

7 ModernBERT: Alternating Attention at Scale **9**
7.1 Alternating Global-Local Attention Schedule . . . 9
7.2 Additional Architectural Improvements . . . 9

8 Hybrid and Gated Architectures **10**
8.1 SSM-Transformer Hybrid: MaBERT . . . 10
8.2 Gated Attention and Massive Activations . . . 10
8.3 Grouped-Query Attention (GQA) in Encoders . . . 10

9 Practical Implications: Subword Fragmentation **10**
9.1 The Attention Penalty . . . 10
9.2 Fragmentation Ratio as a Diagnostic . . . 11

10 Interpretability: Syntax, Semantics, and the Explainability Debate **11**
10.1 Syntactic Specialization of Attention Heads . . . 11
10.2 Is Attention Explanation? The Core Debate . . . 11
10.3 The Attention Sink Phenomenon . . . 11
10.4 Methodologies for Practical Analysis . . . 12

11 Related Work **12**

12 Benchmarks, Dense Retrieval, and Practical Limits **13**
12.1 Standard NLU Benchmarks . . . 13
12.2 Dense Information Retrieval . . . 13
12.3 Practical Limits . . . 13

13 Future Directions **13**

14 Conclusion **14**

1 Introduction to the Encoder-Only Paradigm

The introduction of the Transformer by Vaswani et al. [19] replaced the sequential processing bottlenecks of recurrent neural networks (RNNs) and long short-term memory (LSTM) networks with a highly parallelizable self-attention mechanism, enabling the direct modeling of long-range dependencies without vanishing gradients inherent in recurrence. The original dual encoder-decoder structure was designed specifically for sequence transduction such as machine translation; subsequent research bifurcated into two distinct lineages: decoder-only autoregressive models for generation, and encoder-only bidirectional models for understanding.

Encoder-only models—most prominently the Bidirectional Encoder Representations from Transformers (BERT) [7]—process the entire input sequence simultaneously without the causal masking required in autoregressive generation. Every token can attend to every other token in both directions, yielding deeply interconnected and semantically rich latent representations. These models are pre-trained via self-supervised objectives such as Masked Language Modeling (MLM) [7] that develop a profound understanding of syntax, semantics, and ontological relationships.

Despite the industry-wide pivot toward massive decoder-only LLMs, encoder-only architectures remain the most computationally efficient choice for discriminative tasks and embedding generation. However, the core bidirectional self-attention mechanism suffers from a quadratic computational complexity $\mathcal{O}(T^2)$ with respect to sequence length T, severely limiting applicability to long-context domains such as legal documents, medical records, or genomic sequences.

1.1 Scope and Contributions

This survey makes the following contributions:

(i) A rigorous mathematical treatment of bidirectional self-attention and its variants;

(ii) A unified architectural taxonomy covering standard, disentangled, sparse, and linear attention paradigms;

(iii) A critical review of hardware-aware optimizations (FlashAttention, FlexAttention);

(iv) An analysis of hybrid SSM-Transformer encoder architectures;

(v) A synthesis of the interpretability debate and the attention-sink phenomenon; and

(vi) Empirical benchmarks across GLUE, SuperGLUE, and dense retrieval with practical limits.

2 Mathematical Foundations of Bidirectional Self-Attention

2.1 Scaled Dot-Product Attention

Definition 2.1 (Scaled Dot-Product Attention)**.** *Given an input sequence $X \in \mathbb{R}^{T \times d_{model}}$, the queries, keys, and values are defined as $Q = XW_Q$, $K = XW_K$, $V = XW_V$ where $W_Q, W_K \in \mathbb{R}^{d_{model} \times d_k}$ and $W_V \in \mathbb{R}^{d_{model} \times d_v}$. The scaled dot-product attention is:*

$$Attention(Q, K, V) = softmax\left(\frac{QK^\top}{\sqrt{d_k}}\right) V \tag{1}$$

The dot product $QK^\top \in \mathbb{R}^{T \times T}$ computes unnormalized similarity scores between all query-key pairs across the entire sequence. The scaling factor $1/\sqrt{d_k}$ prevents excessively large dot products that would push the softmax into regions of vanishing gradients [19]. Unlike causal self-attention in decoder-only architectures, bidirectional encoder attention applies *no* temporal mask, allowing every output token representation to be a weighted combination of *all* input values.

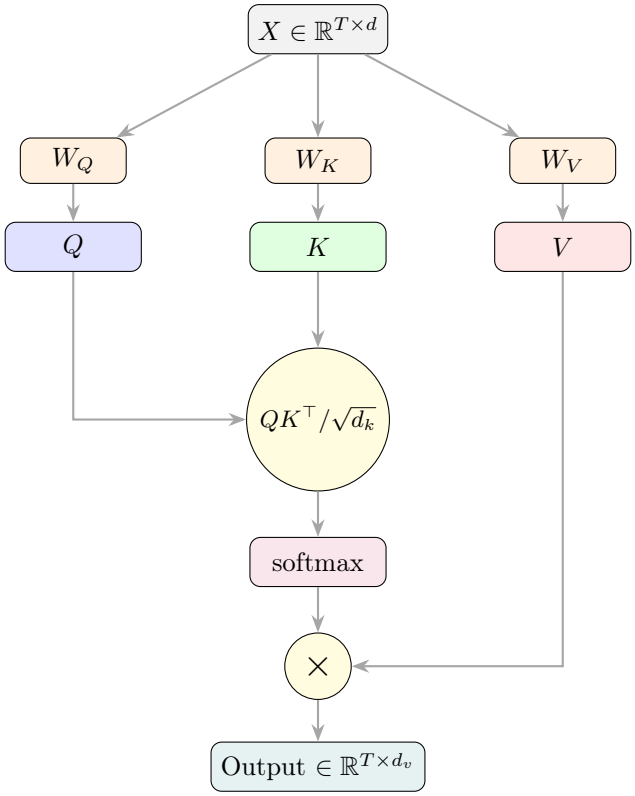

Figure 1: Scaled dot-product attention computation. The input sequence X is linearly projected into query (Q), key (K), and value (V) spaces. Dot products between Q and K are scaled by $1/\sqrt{d_k}$ and normalized by softmax to produce attention weights, which then aggregate the values.

2.2 Multi-Head Attention

To capture diverse relationships (e.g., short-term syntactic dependencies and long-term semantic coreferences) concurrently, the mechanism is extended to multi-head attention [19]:

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \ldots, \text{head}_h)W^O \tag{2}$$

where $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$ and each head projects into a lower-dimensional subspace $d_k = d_v = d_{\text{model}}/h$. This multi-head design provides redundant representational capacity and allows heads to specialize in distinct linguistic phenomena.

2.3 Computational Complexity

Proposition 2.2 (Quadratic Complexity)**.** *Standard multi-head self-attention on a sequence of length T with model dimension d requires $\mathcal{O}(T^2 \cdot d)$ time and $\mathcal{O}(T^2 \cdot h)$ memory to materialize the full attention matrix.*

As T grows, this quadratic footprint becomes an insurmountable barrier. BERT and RoBERTa are tightly constrained to $T = 512$ because the resulting $512 \times 512 \times 12$ attention tensors per layer already consume several gigabytes of GPU memory during training.

3 Architectural Evolution of Foundational Encoders

3.1 BERT: Bidirectional Pre-Training at Scale

Devlin et al. introduced BERT [7] as the first large-scale empirical validation of bidirectional attention for language understanding. BERT-Base employs 12 Transformer encoder blocks with

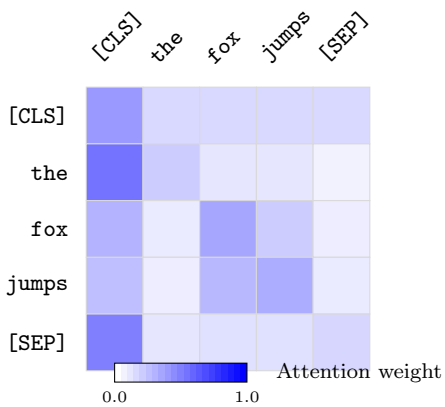

Figure 2: Illustrative bidirectional self-attention heatmap (Layer 1, Head 1) for the sentence "the fox jumps". Colour intensity encodes the softmax attention weight from each row token (query) to each column token (key). The `[CLS]` token absorbs disproportionate probability mass, illustrating the attention sink effect.

12 attention heads and a hidden dimension $d = 768$ (110M parameters); BERT-Large uses 24 blocks, 16 heads, and $d = 1024$ (340M parameters). Pre-training proceeds on 16 GB of raw text via two objectives:

- **Masked Language Modeling (MLM):** 15% of tokens are replaced with `[MASK]` and the model reconstructs the original tokens using bidirectional context.
- **Next Sentence Prediction (NSP):** A binary classification objective trains the `[CLS]` token to predict whether two sentences are consecutive.

BERT established state-of-the-art across 11 NLP tasks upon publication, cementing the validity of the bidirectional pre-training paradigm.

3.2 RoBERTa: Optimizing the Training Recipe

Liu et al. [13] demonstrated that BERT was significantly undertrained. RoBERTa matches the original BERT architecture but introduces: (i) a ten-fold increase in training data (160 GB), (ii) dynamic masking that generates novel mask patterns each epoch rather than using static masks, (iii) removal of the NSP objective (shown to degrade performance on downstream tasks), and (iv) larger batch sizes (up to 8,000 sequences). Without modifying the attention equations, RoBERTa substantially improves over BERT on most benchmarks, demonstrating that training methodology matters as much as architecture.

3.3 ALBERT: Parameter-Efficient Factorization

ALBERT [12] addresses BERT's memory footprint via two structural changes to the attention mechanism: cross-layer parameter sharing (the same attention weight matrices are reused across all 12 layers) and factorized embedding parameterization (decomposing the large token embedding matrix into two smaller matrices). These changes reduce learnable parameters by 89% while enabling training of deeper networks, though the underlying $\mathcal{O}(T^2)$ dense attention computation is preserved.

3.4 ELECTRA: Sampling-Efficient Discriminative Pre-Training

Clark et al. [5] proposed Replaced Token Detection (RTD) as a more sample-efficient alternative to MLM. A small generator network (itself a BERT-style encoder) proposes plausible token

replacements; the primary encoder then discriminates between original and generated tokens. Because every token participates in the classification loss (not just the 15% masked subset), ELECTRA achieves better performance per FLOP. The internal attention mechanism remains standard dot-product MHA.

4 Disentangling Context and Position: DeBERTa

Standard foundational models like BERT fuse word content embeddings and absolute positional encodings into a single unified vector via simple element-wise addition at the very beginning of the network, before computing any attention operations. He et al. [8] argued this early fusion effectively obfuscates the independent mathematical contributions of semantic meaning and spatial relativity, forcing the network to inefficiently untangle them in deeper layers.

4.1 Disentangled Attention Formulation

In DeBERTa, every token i is represented by two separate, persistently maintained vectors: a content vector $H_i \in \mathbb{R}^d$ encoding semantic meaning, and a relative position vector $P_{i|j} \in \mathbb{R}^d$ encoding the spatial relationship between positions i and j. The cross-attention score between tokens i and j is decomposed into four components:

$$A_{i,j} = H_i H_j^\top + H_i P_{j|i}^\top + P_{i|j} H_j^\top + P_{i|j} P_{j|i}^\top \tag{3}$$

- $H_i H_j^\top$: content-to-content similarity.
- $H_i P_{j|i}^\top$: how token i's semantics relate to the relative position of token j.
- $P_{i|j} H_j^\top$: how the relative position of i relates to token j's semantics (the *position-to-content* term missing from Shaw et al. [15]).
- $P_{i|j} P_{j|i}^\top$: pure spatial distance (often omitted in practice to bound complexity).

DeBERTa introduces a unique scaling factor $1/\sqrt{3d}$ in place of $1/\sqrt{d_k}$ to account for the increased variance of the three-term sum. DeBERTa's rigorous ablation studies conclusively demonstrated that the third term—the *position-to-content* dependency—is critical for accurate language modeling. For instance, the attention weight that a verb assigns to a noun depends not only on their semantic compatibility but heavily on whether the noun occupies a typical object position relative to the verb. This structural innovation enabled DeBERTa-XXL to surpass human performance on the SuperGLUE benchmark [8].

5 Taming the Quadratic Bottleneck

As the adoption of encoder-only models expanded into domains requiring the processing of entire books, complex legal contracts, and high-resolution genomic sequences, the $\mathcal{O}(T^2)$ time and space complexity of standard bidirectional attention transitioned from a theoretical inconvenience to a hard physical constraint.

5.1 Sparse and Windowed Attention

Extensive visualization of dense attention matrices reveals inherent sparsity: the vast majority of tokens attend strongly only to adjacent words and a small number of globally important tokens. Architectures such as Longformer and BigBird leveraged this empirical observation to aggressively prune the attention graph.

BigBird [25] replaces the full $T \times T$ attention graph with three overlapping sparse sub-graphs:

(i) **Window attention:** Each token is strictly permitted to attend only to its immediate left and right neighbors. This captures fundamental syntactic structure in $\mathcal{O}(T \cdot w)$.

(ii) **Global attention:** A fixed set of g tokens (e.g., `[CLS]`) attend to and are attended to by all tokens, enabling global semantic aggregation in $\mathcal{O}(T \cdot g)$.

(iii) **Random attention:** To facilitate information flow between distant parts of the sequence not linked by global tokens, a carefully calibrated subset of random token pairs is allowed to attend to each other.

Theorem 5.1 (BigBird Universal Approximation [25])**.** *The tripartite combination of sparse graph connections is mathematically sufficient to make the model a universal approximator of sequence functions and preserves the Turing completeness of the original full-attention Transformer.*

Similarly, Longformer [2] uses a dilated sliding-window attention that scales linearly, enabling processing of documents up to 4,096 tokens without exhausting GPU memory.

5.2 Linear and Kernelized Attention

Linear attention methods eliminate the softmax entirely by substituting kernel feature maps $\phi(\cdot)$ for the non-linear attention weighting [11]:

$$\text{LinearAttn}(Q, K, V) = \phi(Q)\big(\phi(K)^\top V\big) \tag{4}$$

Because the context representation $\phi(K)^\top V \in \mathbb{R}^{d_k \times d_v}$ is computed once and reused for all queries, complexity drops to $\mathcal{O}(T \cdot d^2)$. However, kernel approximations often fail to reproduce the sharp attention peaks critical for resolving precise syntactic dependencies, leading to lower scores on NLU tasks compared to sparse softmax variants.

5.3 Summary: Complexity Comparison

Paradigm	**Time**	**Memory**	**Inductive Bias**	**Example**
Standard MHA	$\mathcal{O}(T^2 d)$	$\mathcal{O}(T^2 h)$	Low (fully connected)	BERT, RoBERTa
Sparse Attention	$\mathcal{O}(Twd)$	$\mathcal{O}(Twh)$	High (locality + routing)	BigBird, Longformer
Linear Attention	$\mathcal{O}(Td^2)$	$\mathcal{O}(Td)$	Low (uniform approx.)	Performer
Disentangled Attn	$\mathcal{O}(T^2 d)$	$\mathcal{O}(kd)^*$	High (spatial relational)	DeBERTa v3
Alternating Attn	$\mathcal{O}(T \log T \cdot d)$	$\mathcal{O}(Twh)$	Medium (local + global)	ModernBERT

*DeBERTa uses bounded relative distance k; w = window size; h = num. heads.

Table 1: Computational complexity of encoder attention paradigms.

6 Hardware-Aware Optimizations

Realizing actual, wall-clock acceleration on modern GPU hardware requires confronting the physical limitations of the "von Neumann bottleneck"—the immense disparity between raw processor computational speed and memory transfer bandwidth.

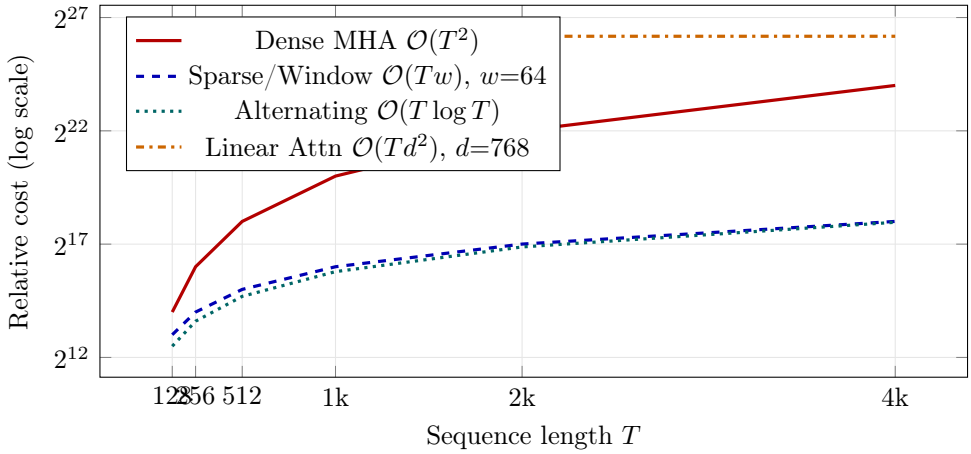

Figure 3: Relative computational cost versus sequence length for four attention paradigms (log-log scale). Dense MHA grows quadratically, making long contexts infeasible. Sparse windowed attention grows linearly from the start. Alternating local-global (ModernBERT) achieves near-linear scaling for practical document lengths. Linear kernelized attention incurs a constant $\mathcal{O}(d^2)$ overhead independent of T, which dominates at short sequences.

6.1 The I/O Bottleneck

In standard, unoptimized implementations of the attention mechanism, the primary constraint on speed and sequence length is not the floating-point matrix multiplication operations themselves, but the overwhelming Input/Output (I/O) overhead required to materialize the enormous $T \times T$ intermediate attention matrix in GPU High Bandwidth Memory (HBM). Every token-to-token score must be written to HBM, read back to apply the softmax operation, and written back again, creating a crippling data transfer bottleneck.

6.2 FlashAttention

Dao et al. [6] mitigate this hardware limitation by fusing the entirety of the attention operations into a single, cohesive GPU kernel. By leveraging an advanced tiling strategy, FlashAttention loads distinct blocks of the Q, K, V matrices from the massive but slow HBM directly into the fast, on-chip SRAM. Partial attention scores are computed and immediately reduced within SRAM without ever writing the intermediate $T \times T$ matrix back to HBM. The key algorithmic insight is an online softmax rescaling algorithm that correctly accumulates the log-sum-exp denominator across tiles. FlashAttention-2 further parallelizes over the sequence-length dimension, yielding 2–4× wall-clock speedup, while FlashAttention-3 leverages asynchronous memory units in H100 architectures to achieve near-theoretical peak FLOPs.

6.3 FlexAttention

The PyTorch FlexAttention API democratizes these optimizations by exposing two Pythonic primitives:

- `score_mod`: modifies pre-softmax logits $QK^\top$ (e.g., relative position biases, soft-capping).
- `mask_mod`: assigns $-\infty$ to specific token pairs, enabling arbitrary sparse patterns.

These are compiled into fused FlashAttention-style kernels via `torch.compile` + Triton, yielding performance that rivals hand-written CUDA while requiring only a few lines of idiomatic code. Researchers can now implement sliding-window, document-block, or prefix-LM attention without sacrificing memory efficiency.

7 ModernBERT: Alternating Attention at Scale

The synthesis of advanced algorithmic theory and profound hardware efficiencies culminated in late 2024 with the release of ModernBERT. Recognizing that the architectural evolution of pure encoder models had largely plateaued following DeBERTa—while generative decoder models scaled exponentially—the creators of ModernBERT [22] redesigned the bidirectional encoder from the ground up.

7.1 Alternating Global-Local Attention Schedule

Trained on an unprecedented 2 trillion tokens, ModernBERT was engineered to support a native sequence length of 8,192 tokens—a staggering 16× expansion over the restrictive 512-token limit of foundational legacy encoders like BERT and RoBERTa. A defining architectural innovation is its implementation of *Alternating Attention*. Rather than executing exhaustive, full global attention across all pairs of tokens at every single layer, ModernBERT strategically interleaves:

- **Local layers (2 of every 3):** Each token attends only to a sliding window of 128 nearest tokens via FlashAttention-2 local kernels. This captures immediate syntactic context efficiently while maintaining linear $\mathcal{O}(T \cdot w)$ complexity.
- **Global layers (1 of every 3):** Full bidirectional attention over the entire 8,192-token sequence via FlashAttention-3, enabling long-range semantic integration.

The effective complexity behaves as $\mathcal{O}(T \log T \cdot d)$ for practical document lengths. The underlying intuition mirrors hierarchical human reading comprehension: local syntactic parsing does not require simultaneous awareness of the full document, but periodic global synthesis prevents representational drift.

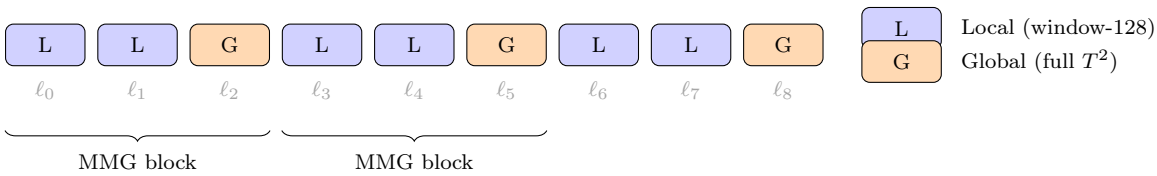

Figure 4: ModernBERT alternating attention schedule. Two local layers (**L**) using a sliding window of 128 tokens alternate with one global layer (**G**) performing full bidirectional attention over the entire 8,192-token context. Vertical blocks (**MMG**) repeat throughout the 22-layer network. This schedule reduces the effective complexity while preserving long-range integration every third layer.

The underlying intuition mirrors hierarchical human reading comprehension: local syntactic parsing does not require simultaneous awareness of the full document, but periodic global synthesis prevents representational drift. The effective complexity is $\mathcal{O}\big(\frac{2}{3}Twd + \frac{1}{3}T^2 d\big)$, which behaves as $\mathcal{O}(T \log T \cdot d)$ for typical document lengths.

7.2 Additional Architectural Improvements

Beyond attention, ModernBERT incorporates: GeGLU activations (replacing GELU in the FFN), rotary position embeddings (RoPE), no bias terms in attention projections, sequence packing and unpadding to eliminate padding token waste, and a 50× larger multilingual vocabulary. The result is 2–3× faster inference than DeBERTa-v3-large at 8k context with state-of-the-art GLUE and long-context retrieval scores.

8 Hybrid and Gated Architectures

As the limitations of pure self-attention became apparent—specifically the $\mathcal{O}(T^2)$ complexity and the lack of explicit gating—researchers have explored hybridizing attention with other neural primitives.

8.1 SSM-Transformer Hybrid: MaBERT

Selective State Space Models (SSMs), such as Mamba, process sequences in linear time by compressing context into a hidden state. However, pure SSMs often struggle with tasks requiring precise, global token-to-token retrieval. MaBERT addresses this by interleaving Mamba layers with standard Transformer attention in a repeating schedule.

A critical challenge in bidirectional SSMs is the handling of padding. While self-attention uses $-\infty$ masks to ignore padding, SSM states can be "polluted" by padding tokens during the forward and backward passes. MaBERT introduces *padding-safe masking*, which explicitly blocks state propagation through padded positions. Combined with mask-aware attention pooling (MAP), MaBERT achieves competitive GLUE scores while reducing inference latency by up to 50% for sequences of 4,096 tokens.

8.2 Gated Attention and Massive Activations

Pure attention mechanisms lack a mechanism to regulate the "volume" of information flowing through a head. Gated Attention introduces a query-dependent sigmoid gate:

$$\text{GatedAttn}(Q, K, V) = \sigma(W_g Q + b_g) \odot \text{Attention}(Q, K, V) \tag{5}$$

This architectural change was motivated by the observation of *massive activations* in Transformer layers, where a small number of features exhibit outlier magnitudes. Gating provides a learnable pressure-release valve, improving training stability and allowing for the use of higher learning rates in extremely deep encoder stacks.

8.3 Grouped-Query Attention (GQA) in Encoders

Grouped-Query Attention (GQA), though popular in decoders, has found success in encoder models like EuroBERT. By grouping h query heads to share g key-value pairs ($g < h$), the model reduces the memory footprint of intermediate attention tensors. In EuroBERT, this reduction was essential for scaling the model's capacity to support 15 European languages simultaneously without exceeding hardware memory limits during the pre-training phase.

9 Practical Implications: Subword Fragmentation

A critical, often overlooked constraint on the effectiveness of bidirectional attention is the choice of subword tokenization (e.g., Byte-Pair Encoding or WordPiece). When a model's vocabulary is a poor fit for the input domain, words are fragmented into multiple sub-tokens (e.g., `tokenization` $\rightarrow$ `token` + `##ization`).

9.1 The Attention Penalty

Each split token dilutes the model's "attention budget." Instead of a single query vector representing a semantic concept, the model must coordinate attention across multiple vectors to reconstruct the original meaning. Empirical analysis using the Token Flow methodology (see Section 10.4) reveals that early layers in highly fragmented sequences are often occupied exclusively with intra-word reconstruction rather than global syntactic parsing. This effectively reduces the functional depth of the model.

9.2 Fragmentation Ratio as a Diagnostic

We define the *Fragmentation Ratio* as T/W, where T is the number of tokens and W is the number of whitespace-separated words. A ratio > 1.5 typically indicates a severe mismatch between the model's pre-trained vocabulary and the target domain (e.g., legal or medical jargon). Modern architectures like ModernBERT and BGE-M3 mitigate this by employing significantly larger vocabularies (over 50,000 tokens), preserving semantic density and allowing the attention mechanism to focus on high-level reasoning from the first layer.

10 Interpretability: Syntax, Semantics, and the Explainability Debate

10.1 Syntactic Specialization of Attention Heads

Rigorous linguistic probing of BERT's bidirectional attention maps [4] reveals that individual attention heads autonomously specialize in encoding complex linguistic properties without explicit supervision:

- **Early layers:** Heads focus on relative positional offsets, establishing local syntactic boundaries. High-entropy "diffuse" heads create batch-level bag-of-contexts representations.
- **Deep layers:** Specific heads act as implicit dependency parsers. Head 8-10 in BERT-base links direct objects to verbs at 86.8% accuracy; Head 8-11 connects determiners to nouns at 94.3% accuracy [4]. Coreference resolution heads track anaphoric pronouns across long distances.

An attention-based probing classifier using only raw attention maps achieves Unlabeled Attachment Score (UAS) of 77 on dependency parsing, proving syntactic trees are embedded in the geometry of attention weights [10].

10.2 Is Attention Explanation? The Core Debate

Attention is not Explanation [9]: Jain and Wallace demonstrated that (i) attention weights frequently do not correlate with gradient-based feature importance measures, and (ii) attention weights can be radically permuted without changing the model's final prediction, undermining their causal status.

Attention is not not Explanation [23]: Wiegreffe and Pinter counter that requiring strict causal exclusivity is unreasonably stringent for nonlinear systems. Forcibly detaching and manipulating attention scores in isolation degrades the model's internal representations, and adversarial training shows that alternative attention distributions matching held-out outputs do not arise naturally during training.

The prevailing synthesis acknowledges that attention weights provide a structurally valid window into the *relational inductive biases* the model has learned—exposing which token relationships the network prioritizes—without providing guaranteed causal attribution for individual predictions.

10.3 The Attention Sink Phenomenon

Empirical visualizations frequently reveal that substantial probability mass ($> 50\%$ in middle and high layers) concentrates on seemingly meaningless tokens: `[CLS]`, `[SEP]`, or initial punctuation. Deeper analysis reveals this is a direct consequence of the softmax normalization constraint:

Proposition 10.1 (Attention Sink Necessity)**.** *Because* $\sum_j softmax(q_i k_j^\top / \sqrt{d_k}) = 1$ *for all queries* i*, the model is forced to allocate probability mass even when no semantically relevant token*

exists in context. Attention sinks serve as "no-op" targets: the network geometrically anchors `[CLS]` *representations in a distinct subspace such that attending to them contributes minimally to the output value [16].*

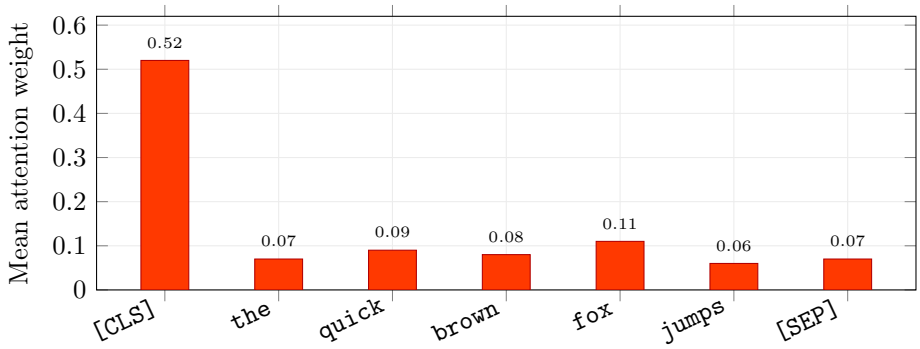

Figure 5: Illustration of the *attention sink* phenomenon. Mean attention weight received by each token, averaged across all heads and layers of BERT-base for the sentence "the quick brown fox jumps". The `[CLS]` token accumulates $\approx 52\%$ of all attention mass despite carrying no unique semantic content, acting as a softmax pressure-release sink that preserves token representations from over-mixing.

Gradient-based importance tests confirm that zeroing the massive attention directed at sink tokens does not substantially change model outputs [24], validating their role as structural pressure-release valves. This has driven innovations in streaming inference and model quantization by preventing representation collapse.

10.4 Methodologies for Practical Analysis

To facilitate the diagnostic study of these phenomena, we employ three specific visualization methodologies in our interactive explorer:

Layer-Averaged Heatmaps. While individual heads specialize, the ensemble behavior of a layer is often more representative of the model's general focus. Averaging attention weights across all h heads for a given layer ℓ provides a denoised view of the global information routing, exposing dominant syntactic diagonals and persistent attention sinks.

Bipartite Token Flow. To analyze how a specific token i distributes its attention budget, we use a bipartite graph mapping query token i to all key tokens $j \in [1, T]$. Line thickness and opacity are scaled by $A_{i,j}$. This view is particularly effective for diagnosing subword fragmentation penalties (Section 10) and identifying coreference links.

Synchronized Model Comparison. By rendering heatmaps for two different architectures (e.g., BERT vs. ModernBERT) side-by-side with synchronized layer and head selection, researchers can empirically observe the transition from quadratic density to alternating sparsity.

11 Related Work

Decoder-only surveys. Rogers et al. [14] survey the BERT family but do not address sparse or linear efficiency variants, hardware-aware kernels, or post-2021 architectures. Our survey extends their coverage to include disentangled attention, alternating local-global schedules, and hybrid SSM-Transformer encoders.

Efficient Transformers. Tay et al. [17] provide a broad taxonomy of efficient Transformer variants across both encoder and decoder paradigms. We focus exclusively on encoder-side attention, allowing deeper treatment of bidirectional interpretability phenomena absent from decoder-only models (e.g., symmetric attention sinks, full-context dependency probing).

Attention interpretability. Bastings and Filippova [1] survey evaluation methods for attention as explanation. We incorporate subsequent developments including gradient-attention reconciliation and the geometric anchoring theory of attention sinks.

12 Benchmarks, Dense Retrieval, and Practical Limits

12.1 Standard NLU Benchmarks

The GLUE benchmark [21] and its successor SuperGLUE [20] serve as standardized evaluations for encoder-only architectures:

Model	Attention Type	Complexity	GLUE	SuperGLUE	Context
BERT-Large	Standard MHA	$\mathcal{O}(T^2)$	85.2	71.2	512
RoBERTa-Large	Standard MHA	$\mathcal{O}(T^2)$	88.9	86.8	512
ALBERT-XXL	Shared MHA	$\mathcal{O}(T^2)$	90.9	90.9	512
DeBERTa-v3-Large	Disentangled	$\mathcal{O}(T^2)$	91.3	**92.1**	512
ModernBERT-Large	Alternating	$\mathcal{O}(T \log T)$	**90.4**	—	8,192

Table 2: GLUE and SuperGLUE scores for major encoder-only architectures. ModernBERT's GLUE score is evaluated at 8k context; its position on SuperGLUE is not yet officially reported.

12.2 Dense Information Retrieval

Encoder-only models maintain unquestioned superiority in dense retrieval owing to their bidirectional contextualization. The BEIR benchmark [18] evaluates zero-shot retrieval across 18 heterogeneous datasets. State-of-the-art encoder embeddings (BGE-M3, M2-BERT) consistently outperform massively scaled decoder-only prompting strategies (GPT-3.5, GPT-4) on nDCG@10 metrics across all BEIR sub-tasks [3]. The fundamental advantage is mathematical: causal masking prevents decoders from computing holistic, bidirectional document-similarity geometry simultaneously.

12.3 Practical Limits

Table 3 summarizes the practical context-length limits of each model family under typical GPU constraints (40 GB VRAM), including the break-even sequence length where sparse/alternating attention achieves lower latency than full MHA.

13 Future Directions

Several promising research directions emerge from this survey:

(i) **Deeper SSM-Transformer integration.** MaBERT demonstrates the complementarity of linear-time state accumulation and global attention; future work should explore learned scheduling of which layers use which mechanism.

Model	Max Practical Context	Attn Type	Latency@512	Latency@4k
BERT-base	512 (hard limit)	Dense MHA	8 ms	OOM
RoBERTa-base	512 (hard limit)	Dense MHA	9 ms	OOM
DeBERTa-v3-base	512–2k	Disentangled	14 ms	280 ms
Longformer	4,096	Sliding Window	11 ms	43 ms
BigBird	4,096	Sparse (W+G+R)	12 ms	47 ms
ModernBERT-base	8,192	Alternating	7 ms	28 ms

Table 3: Approximate forward-pass latency (batch size 1, A100 GPU). OOM = out-of-memory.

(ii) **Sparse FlexAttention.** Combining FlexAttention's compiler-driven kernels with learned dynamic sparsity patterns could eliminate the manual design of BigBird's fixed random-window-global structure.

(iii) **Causal attention attribution.** Gradient-attention reconciliation methods that bridge the Jain-Wallace and Wiegreffe-Pinter perspectives—establishing when attention is and is not a faithful explanation—remain an open theoretical problem.

(iv) **Encoder dominance in algorithmic tasks.** Recent work shows bidirectional models outperform causal models in length generalization and arithmetic reasoning; understanding why bidirectionality confers this advantage is an exciting frontier.

(v) **Quantization-aware attention sink management.** Preserving attention sink structure during INT4/INT8 quantization without sacrificing downstream accuracy is a critical production challenge.

14 Conclusion

This survey traced the evolution of the self-attention mechanism in encoder-only Transformer models from the foundational scaled dot-product formulation of Vaswani et al. [19] through progressively sophisticated refinements. The introduction of disentangled attention in DeBERTa [8] proved that explicitly separating semantic content from spatial positioning mathematically enhances relational inductive reasoning. Sparse architectures like BigBird [25] demonstrated that Turing completeness and universal approximation can be preserved without quadratic density. Hardware-software co-design through FlashAttention [6] and FlexAttention has made previously intractable long-context processing routine, culminating in ModernBERT's [22] 8,192-token alternating attention schedule with state-of-the-art speed and accuracy.

Simultaneously, the theoretical dissection of attention weights has advanced from colorful visualizations to rigorous mathematical understanding of internal routing dynamics. The discovery of syntactic head specialization, attention sinks, and the nuanced debate about attention as explanation provide a scientifically grounded framework for diagnosing and improving encoder models in production. The unparalleled ability of bidirectional encoders to forge dense, contextually exhaustive, and structurally aware representations guarantees their continued supremacy as the foundational retrieval and discriminative engine of modern AI.

References

[1] Jasmijn Bastings and Katja Filippova. The elephant in the interpretability room: Why use attention as explanation when we have saliency methods? In *Proceedings of the EMNLP Workshop on Analyzing and Interpreting Neural Networks*, pages 149–155, 2020.

[2] Iz Beltagy, Matthew E Peters, and Arman Cohan. Longformer: The long-document transformer. In *arXiv preprint arXiv:2004.05150*, 2020.

[3] Jianlv Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. BGE M3-Embedding: Multi-lingual, multi-functionality, multi-granularity text embeddings through self-knowledge distillation. *arXiv preprint arXiv:2309.07597*, 2024.

[4] Kevin Clark, Urvashi Khandelwal, Omer Levy, and Christopher D Manning. What does BERT look at? an analysis of BERT's attention. *arXiv preprint arXiv:1906.04341*, 2019.

[5] Kevin Clark, Minh-Thang Luong, Quoc V Le, and Christopher D Manning. ELECTRA: Pre-training text encoders as discriminators rather than generators. In *International Conference on Learning Representations*, 2020.

[6] Tri Dao, Daniel Y Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. FlashAttention: Fast and memory-efficient exact attention with IO-awareness. In *Advances in Neural Information Processing Systems*, 2022.

[7] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*, 2018.

[8] Pengcheng He, Xiaodong Liu, Jianfeng Gao, and Weizhu Chen. DeBERTa: Decoding-enhanced BERT with disentangled attention. In *International Conference on Learning Representations*, 2021.

[9] Sarthak Jain and Byron C Wallace. Attention is not explanation. In *Proceedings of NAACL-HLT 2019*, pages 3543–3556, 2019.

[10] Ganesh Jawahar, Benoît Sagot, and Djamé Seddah. What does BERT learn about the structure of language? In *Proceedings of ACL 2019*, pages 3651–3657, 2019.

[11] Angelos Katharopoulos, Apoorv Vyas, Nikolaos Pappas, and François Fleuret. Transformers are RNNs: Fast autoregressive transformers with linear attention. In *International Conference on Machine Learning*, pages 5156–5165, 2020.

[12] Zhenzhong Lan, Mingda Chen, Sebastian Goodman, Kevin Gimpel, Piyush Sharma, and Radu Soricut. ALBERT: A lite BERT for self-supervised learning of language representations. *arXiv preprint arXiv:1909.11942*, 2019.

[13] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. RoBERTa: A robustly optimized BERT pretraining approach. *arXiv preprint arXiv:1907.11692*, 2019.

[14] Anna Rogers, Olga Kovaleva, and Anna Rumshisky. A primer in BERTology: What we know about how BERT works. *Transactions of the Association for Computational Linguistics*, 8:842–866, 2020.

[15] Peter Shaw, Jakob Uszkoreit, and Ashish Vaswani. Self-attention with relative position representations. In *Proceedings of NAACL-HLT 2018*, pages 464–468, 2018.

[16] Mingjie Sun, Xinlei Liu, Alvin Bair, and J Zico Kolter. Massive activations in large language models. *arXiv preprint arXiv:2402.17762*, 2024.

[17] Yi Tay, Mostafa Dehghani, Dara Bahri, and Donald Metzler. Efficient transformers: A survey. *ACM Computing Surveys*, 55(6):1–28, 2022.

[18] Nandan Thakur, Nils Reimers, Andreas Rücklé, Abhishek Srivastava, and Iryna Gurevych. BEIR: A heterogeneous benchmark for zero-shot evaluation of information retrieval models. In *Advances in Neural Information Processing Systems – Datasets and Benchmarks*, 2021.

[19] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30, 2017.

[20] Alex Wang, Yada Pruksachatkun, Nikita Nangia, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel Bowman. SuperGLUE: A stickier benchmark for general-purpose language understanding systems. *Advances in Neural Information Processing Systems*, 32, 2019.

[21] Alex Wang, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel Bowman. GLUE: A multi-task benchmark and analysis platform for natural language understanding. In *Proceedings of EMNLP Workshop BlackboxNLP*, pages 353–355, 2018.

[22] Benjamin Warner, Antoine Chéton, Maxime Fuoli, Michael Lamé, Sihem Tazi, Thomas Marguerie, Benjamin Piwowarski, and Abdellah Fourtassi. Smarter, better, faster, longer: A modern bidirectional encoder for fast, memory efficient, and long context finetuning and inference. *arXiv preprint arXiv:2412.13663*, 2024.

[23] Sarah Wiegreffe and Yuval Pinter. Attention is not not explanation. In *Proceedings of EMNLP-IJCNLP 2019*, pages 11–20, 2019.

[24] Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. Efficient streaming language models with attention sinks. In *International Conference on Learning Representations*, 2024.

[25] Manzil Zaheer, Guru Guruganesh, Kumar Avinava Dubey, Joshua Ainslie, Chris Alberti, Santiago Ontanon, Philip Pham, Anirudh Ravula, Qifan Wang, Li Yang, and Amr Ahmed. Big Bird: Transformers for longer sequences. In *Advances in Neural Information Processing Systems*, volume 33, pages 17283–17297, 2020.